

▼ Introduction and Problem Definition

Background

Bike sharing systems are new generation of traditional bike rentals where whole process from membership, rental and return back has become automatic. Through these systems, user is able to easily rent a bike from a particular position and return back at another position. Currently, there are about over 500 bike-sharing programs around the world which is composed of over 500 thousands bicycles. Today, there exists great interest in these systems due to their important role in traffic, environmental and health issues.

Apart from interesting real world applications of bike sharing systems, the characteristics of data being generated by these systems make them attractive for the research. Opposed to other transport services such as bus or subway, the duration of travel, departure and arrival position is explicitly recorded in these systems. This feature turns bike sharing system into a virtual sensor network that can be used for sensing mobility in the city. Hence, it is expected that most of important events in the city could be detected via monitoring these data.

Data Set

Bike-sharing rental process is highly correlated to the environmental and seasonal settings. For instance, weather conditions, precipitation, day of week, season, hour of the day, etc. can affect the rental behaviors. The core data set is related to the two-year historical log corresponding to years 2011 and 2012 from Capital Bikeshare system, Washington D.C., USA. In summary, in this project, I will be working with the first six steps in the ML pipeline.

Saved successfully!

A problem is said to be regression problem when the output of the variable gives a real or continuous value. Since the Target in the Dataset bike-dataset hour is 'cnt', and it gives a real value. Then the models to be used is the linear regression.

problem statement

The Bike -sharing dataset was obtained so that it can be used in ML model training to build a predictive model that answers the question: "what is the bike rental count hourly, what are the environmental conditions such as the environmental and seasonal settings"? given their data.

▼ Data Ingestion

Before working with the data I need to import all relevant Python modules. Also, I'll read in the dataset to be used "bike-dataset hour".

Then I'll check the first ten rows of the dataframe to have a glimpse of what is in the data

i'll further check which data is categorical or numerical.

```
import numpy as np
import pandas as pd
import sklearn.model_selection
import sklearn.preprocessing
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/bike-dataset hour.csv")
df.head(10)
```

	instant	dteday	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered
0	1	2011-01-01	1	0	1	0	0	6	No	1	0.24	0.2879	0.81	0.0000	3	13
1	2	2011-01-01	1	0	1	1	0	6	No	1	0.22	0.2727	0.80	0.0000	8	32
2	3	2011-01-01	1	0	1	2	0	6	No	1	0.22	0.2727	0.80	0.0000	5	27
3	4	2011-01-01	1	0	1	3	0	6	No	1	0.24	0.2879	0.75	0.0000	3	10

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   instant     17379 non-null  int64
1   dteday      17379 non-null  object
2   season      17379 non-null  int64
3   yr          17379 non-null  int64
4   mnth        17379 non-null  int64
5   hr          17379 non-null  int64
6   holiday     17379 non-null  int64
7   weekday     17379 non-null  int64
8   workingday  15595 non-null  object
9   weathersit   17379 non-null  int64
10  temp        15595 non-null  float64
11  atemp       15595 non-null  float64
12  hum         17379 non-null  float64
13  windspeed   17379 non-null  float64
14  casual      17379 non-null  int64
15  registered  17379 non-null  int64
16  cnt         17379 non-null  int64
dtypes: float64(4), int64(11), object(2)
memory usage: 2.3+ MB
```

df.describe()

	instant	season	yr	mnth	hr	holiday	weekday	weathersit	temp
count	17379.00000	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000	15595.000000
mean	8690.0000	2.501640	0.502561	6.537775	11.546752	0.028770	3.003683	1.425283	0.496451
std	5017.0295	1.106918	0.500008	3.438776	6.914405	0.167165	2.005771	0.639357	0.192580
min	1.0000	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000	1.000000	0.020000
25%	4345.5000	2.000000	0.000000	4.000000	6.000000	0.000000	1.000000	1.000000	0.340000
50%	8690.0000	3.000000	1.000000	7.000000	12.000000	0.000000	3.000000	1.000000	0.500000
75%	13034.5000	3.000000	1.000000	10.000000	18.000000	0.000000	5.000000	2.000000	0.660000
max	17379.0000	4.000000	1.000000	12.000000	23.000000	1.000000	6.000000	4.000000	1.000000

```
#we'll extract the categorical features
categorical = [var for var in df.columns if df[var].dtype=='O']

print('There are {} categorical variabes \n'.format(len(categorical)))

print('They are: ', categorical)

There are 2 categorical variabes

They are:  ['dteday', 'workingday']

#we'll extract the numerical features
numerical = [var for var in df.columns if df[var].dtype!='O']

print('There are {} numerical variabes \n'.format(len(numerical)))

print('They are: ', numerical)

There are 15 numerical variabes

They are:  ['instant', 'season', 'yr', 'mnth', 'hr', 'holiday', 'weekday', 'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'casu
```

Findings from data ingestion

From the above, it is observed that the bike-dataset has a total 17379 rows and 17 columns. It has two Categorical data and 15 Numerical data.

▼ Data Preparation

It is assumed that

- Missing values are in the 'temp' and 'atemp' columns.
- The peak usage hours are: 7-9AM and 4-7PM on working days, and 10am-4pm on non-working days.
- At night (10pm-4am) the bike rentals are low
- If the humidity or wind-speed is high, the number of rentals decreases.

In this case, based on the Exploratory data analysis that will be performed, I want to do the following:

Fill the missing values in the temperature columns automatically with values that would most closely mirror the actual temperature.

- Create a new field that indicates whether it is a peak time or not
- Create a new field that indicates whether it is night time or not
- Remove all fields containing information about specific dates ('yr', 'mnth', 'dteday'), 'casual' and 'registered' and any other variables

Saved successfully!

▼ 3.1 Handling missing values

Based on the information above, missing values are in the temperature column.

Replacing the missing values for the temperature column automatically with values that would most closely mirror the actual temperature.

```
# median
df['temp'].fillna(value = df.temp.median(), inplace = True)

#median
df['atemp'].fillna(value = df.atemp.median(), inplace = True)
```

- ▶ Handling the hr column and creating a new column called peak_period and night_time. Here, I utilized **binning** as it was more convenient to use.

[] ↳ 8 cells hidden

▶ 3.2 Dropping columns that are not needed

The following columns are no longer needed in the dataframe

1. 'hr'
2. 'year'
3. 'mnth'
4. 'dteday'
5. 'casual'
6. 'registered'

[] ↳ 3 cells hidden

▶ 3.3 Encode the data

Encoding the data is the last step in the data preparation stage. Here, we are going to transform non-numerical values into numerical ones so that our models can, later on, learn from that data.

The ordinal data that I have prepared and does not need to be changed further is the variables *night_time*. *Holiday* Column does not need to be encoded too because it is already encoded.

Label Encoding: The variable that needs to be label encoded is working day and peak_period. I would simply turn the various categorical values into numeric values by mapping Yes-->1 and No-->0.

One Hot Encoding: The other remaining variables *Season*, *Weather_Sit*, and others are not ordinal and have each different categories. Since the order of those categories does not matter and we do not have too many features, we can use One-Hot Encoding. For this, each unique value for each variable gets its own column.

[] 2 cells hidden

4) Data Segregation

Now that the data has been prepared, it should be segregated into a train/test or train/val/test split. The training set is used to develop and train my model, the test set is used to obtain the performance of the model.

Since our dataset is yet to be splitted into test and train data already, I am going to split it the data by electing the first 80 percent and save them in a dataframe variable *train_df* and the latter 20% in a variable *test_df*. Then print out the first 10 lines of each subset. I will be doing a random or automatic selection as it is usually better to calculate the results automatically using percentages rather than manually selecting the top 80%.

```
# splitting the dataset into 80% train and 20% test.
train_df = df[:int(0.8*df.shape[0])]
test_df = df[int(0.8*df.shape[0]):]

print(train_df.head(10))
print(test_df.head(10))
```

Saved successfully!

```
as above and save them in variables X and y

y = df['cnt']
#shuffle and extract an 80/20 split of the data in the dataframe df using sklearn.model_selection.train_test_split
X_train, X_test, y_train, y_test = sklearn.model_selection.train_test_split(X, y, test_size=0.2, random_state=42)
# Printing out the top few lines of each dataframe to check that the split was correct
print(X_train.head())
print(X_test.head())
print(y_train.head())
print(y_test.head())

#feature scaling

sc = StandardScaler()
X = sc.fit_transform(X)
X
```

Model training 1

```
#using linear regression
model = LinearRegression()

#fitting the linear regression model
model.fit(X_train, y_train)

LinearRegression()

training_score = model.score(X_train, y_train)
test_score = model.score(X_test, y_test)

print(f'Linear Regression -Training set score: {training_score:.2f}')
print(f'Linear Regression -Test set score: {test_score:.2f}')

Linear Regression -Training set score: 0.60
Linear Regression -Test set score: 0.59

prediction = model.predict(X_test)

print('The R^2 score : ')

r2_score(y_test,prediction)

The R^2 score :
0.5875062516813065
```

The R^2 score indicate that the score isn't good enough for modelling so I will be retrain the model using Lasso and Ridge

```
mean_squared_error(y_test, prediction)
```

```
13061.78421422828
```

```
lasso = Lasso()
ridge = Ridge()
```

```
lasso.fit(X_train, y_train)
ridge.fit(X_train, y_train)
```

```
Ridge()
```

```
y_prediction_o = model.predict(X_test)
y_lasso_o = lasso.predict(X_test)
y_ridge_o = ridge.predict(X_test)
```

```
print("Model(no outliers)\t RMSE \t\t MSE \t\t MAE \t\t R2 ")
print("""LinearRegression \t{:.2f} \t\t{:.2f} \t\t{:.2f} \t\t{:.2f}""".format(mean_squared_error(y_test, prediction, squared=False), mean_s
print("""LassoRegression \t{:.2f} \t\t{:.2f} \t\t{:.2f} \t\t{:.2f}""".format(mean_squared_error(y_test, y_lasso_o, squared=False), mean_s
print("""RidgeRegression \t{:.2f} \t\t{:.2f} \t\t{:.2f} \t\t{:.2f}""".format(mean_squared_error(y_test, y_ridge_o, squared=False), mean_s
```

Saved successfully!



		MSE	MAE	R2	
LinearRegression	114.23	13061.78	85.59	0.59	
LassoRegression	114.83	13185.89	85.78	0.58	
RidgeRegression	114.27	13058.38	85.59	0.59	

```
rtraining_score = ridge.score(X_train, y_train)
rtest_score = ridge.score(X_test, y_test)
```

```
print(f'Ridge Regression -Initial Training set score: {rtraining_score:.2f}')
print(f'Ridge Regression -Initial Test set score: {rtest_score:.2f}')
```

```
Ridge Regression -Initial Training set score: 0.60
Ridge Regression -Initial Test set score: 0.59
```

After retraining, the score remains the same. So I will try to vary the value of alpha to see any effect

```
ridge = Ridge(alpha= 0.4).fit(X_train, y_train)
```

```
rtraining_score = ridge.score(X_train, y_train)
rtest_score = ridge.score(X_test, y_test)
```

```
print(f'Ridge Regression -Training set score: {rtraining_score:.2f}')
print(f'Ridge Regression -Test set score: {rtest_score:.2f}')
```

```
Ridge Regression -Training set score: 0.60
Ridge Regression -Test set score: 0.59
```

After changing alpha severally, the score remains the same, so I am going to leave it at that

```
model.coef_
```

```
array([-22.48387986,  1.87992361,  4.56556295, 199.87795931,
        113.34805398, -190.48350648, -27.00012952, 209.94380339,
        -44.62621084, -32.65060931,  4.81534393, -9.62027468,
         37.45554006,  30.66314708,  31.26701334, -5.89736124,
        -56.03279918])
```

```
model.intercept_
```

```
80.51879918572621
```

▼ Model Training2

using regression tree

```
# importing the regressor
from sklearn.tree import DecisionTreeRegressor
```

```
# create a regressor object
```

```
regressor = DecisionTreeRegressor(random_state = 42)
```

```
# fit the regressor with X and Y data
```

```
regressor.fit(X, y)
```

```
DecisionTreeRegressor(random_state=0)
```

```
# predicting a new value
```

```
regressor.fit(X_train, y_train)
```

```
predictions = regressor.predict(X_test)
```

```
predictions
```

```
array([451., 13., 10., ..., 62., 341., 372.])
```

X_test

	holiday	weekday	workingday	temp	atemp	hum	windspeed	peak_period	night_time	season_1	season_2	season_3	season_4
12830	0	6	0	0.80	0.6970	0.27	0.1940	0	0	0	0	1	0
8688	1	1	0	0.50	0.4848	0.41	0.2239	0	1	1	0	0	0
7091	0	5	1	0.32	0.3030	0.66	0.2836	0	1	0	0	0	1
			1	0.78	0.7121	0.52	0.3582	1	0	0	1	0	0
431	0	4	1	0.26	0.2273	0.56	0.3881	0	1	1	0	0	0
...
6759	0	5	1	0.56	0.5303	0.94	0.1642	0	0	0	0	0	1
13989	0	6	0	0.64	0.5909	0.78	0.1940	0	1	0	0	1	0
173	0	6	0	0.20	0.1818	0.59	0.3582	1	0	1	0	0	0
16192	1	1	0	0.48	0.4697	0.77	0.1642	1	0	0	0	0	1
8211	0	2	1	0.32	0.3182	0.49	0.1642	1	0	0	0	0	1

3476 rows × 17 columns

y_test

```
12830    425
8688      88
7091       4
12230    526
431       13
...
6759      17
13989     85
173       98
16192    266
8211     267
Name: cnt, Length: 3476, dtype: int64
```

FROM X_test, above, we can see that our model predicted the first cnt of bike rented in X_test (row 12830) to be 425 bikes. The second possum 8688) is estimated to be only 88 bikes.

▼ Checking the performance of the model with the root mean square error (RMSE)

```
tree_rmse = np.sqrt(mean_squared_error(y_test, predictions))
tree_rmse
```

```
139.6397060105962
```

Conclusion

After using linear regression to train the data, it was noticed that the r^2 score was not good enough to achieve a good model. The RMSE score was also too large. The data was also regularized using Lasso and Ridge to see if shrinkage will have effect on the model but it was noticed that it had no effect on it.

Using decion tree model, the regression tree was also used to model the data. The RMSE for the linear regression model is 13061.78 (rounded), (which is very outrageous) while the RMSE for the tree regression model is 139.64 (rounded). The tree regression model performs almost 10 times better than the linear regression model. This shows a profound difference in the two model. Thus it is concluded that the tree model is batter than the linear regression model.

▼ References

- Bike-sharing dataset : <http://capitalbikeshare.com/system-data>
- Extracting categorical and numerical data: <https://stackoverflow.com/questions/26924904/check-if-dataframe-column-is-categorical>
- Handling missing values: Class notes
- Binning of the peak_period and night_time: Class note
- Dropping columns: <https://www.geeksforgeeks.org/how-to-drop-one-or-multiple-columns-in-pandas-dataframe/>
- Data Encoding: Class note
- Data Segregation : Class notes
- Feature Scaling: <https://www.geeksforgeeks.org/ml-feature-scaling-part-2/>
- Deciding the best Metric: <https://medium.com/analytics-vidhya/mae-mse-rmse-coefficient-of-determination-adjusted-r-squared-which-metric-is-bet>
- Lasso and Ridge for regularization: <https://www.geeksforgeeks.org/implementation-of-lasso-ridge-and-elastic-net/>
- Regression Tree: <https://data36.com/regression-tree-python-scikit-learn/>
- Regression Tree: <https://www.geeksforgeeks.org/> and [decision-tree-implementation-python/](#) , [-decision-tree-regression-using-sklearn/](#)

Saved successfully!

